

PYTHON NOTES

with internals

Yang Bo

RDBJ
Admaster

2015.5.25

Outline of the talk

- Python Object
- Built-in Types
- Memory Management
- Reference Count and GC
- ENV: Python in this slides means CPython2.7.6

Python Object

```
#define PyObject_HEAD          \
    _PyObject_HEAD_EXTRA      \
    Py_ssize_t ob_refcnt;      \
    struct _typeobject *ob_type;

typedef struct _object {
    PyObject_HEAD
} PyObject;

#define PyObject_VAR_HEAD      \
    PyObject_HEAD              \
    Py_ssize_t ob_size; /* Number of items in variable part */

typedef struct {
    PyObject_VAR_HEAD
} PyVarObject;
```

object.h

Python Object

- Foundations of Everything(internally)
- Not change size once created
- PyObject / PyVarObject
- Reference Count
- Type Info: `_typeobject`
- Size Info
- Layout

Python Object

- inheritance
- polymorphism using PyObject
- use PyObject * not PyXXObject * when communicating

```
long PyObject_Hash(PyObject *v) {  
    PyTypeObject *tp = v->ob_type;  
    if (tp->tp_hash != NULL)  
        return (*tp->tp_hash)(v);  
    .....  
}
```

object.c

Python Object

```
#define _Py_Dealloc(op) ( \
    _Py_INC_TPFREES(op) _Py_COUNT_ALLOCS_COMMA \
    (*Py_TYPE(op)->tp_dealloc)((PyObject*)(op))) \
#endif /* !Py_TRACE_REFS */

#define Py_INCREF(op) ( \
    _Py_INC_REFTOTAL _Py_REF_DEBUG_COMMA \
    ((PyObject*)(op))->ob_refcnt++) \

#define Py_DECREF(op) \
do { \
    if (_Py_DEC_REFTOTAL _Py_REF_DEBUG_COMMA \
        --((PyObject*)(op))->ob_refcnt != 0) \
        _Py_CHECK_REFCNT(op) \
    else \
        _Py_Dealloc((PyObject*)(op)); \
} while (0)
```

object.h

Python Object/object and type

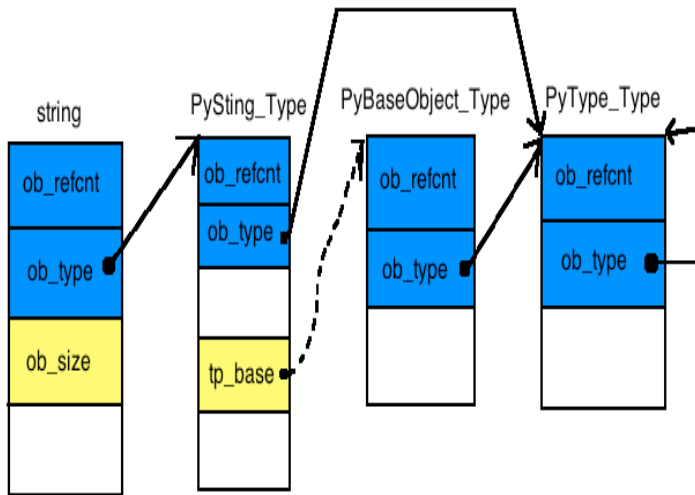
<type 'object'> and <type 'type'>

- 'object' in python is PyObject_Type indeed, not PyObject
- 'object' is base class of everything, including 'type'
- type of all types is 'type'
- type of 'object' is 'type'

```
PyAPI_DATA(PyTypeObject) PyType_Type; /* built-in 'type' */  
PyAPI_DATA(PyTypeObject) PyObject_Type; /* built-in 'object'  
    ' */
```

object.h

Python Object/example



Built-in Types/Integer

```
typedef struct {  
    PyObject_HEAD  
    long ob_ival;  
} PyIntObject;
```

intobject.h

Built-in Types/Integer

- cache small: [-5, 257)
- cache block: PyIntBlock, each 82
- block_list and free_list
- gc: int_dealloc
- memory usage: largest quantity of integers at one exact time, not how many integers totally had in history

Built-in Types/Integer

```
In [6]: %timeit float("1233")
```

The slowest run took 13.29 times longer than the fastest. This could mean that an intermediate result `is` being cached
1000000 loops, best of 3: 233 ns per loop

```
In [7]: %timeit int("1233")
```

The slowest run took 5.72 times longer than the fastest. This could mean that an intermediate result `is` being cached
1000000 loops, best of 3: 708 ns per loop

Why? Let's observe `int_new(objects/intobjects.c)` and `float_new(objects/floatobjects.c)`.

Built-in Types/String

```
typedef struct {
    PyObject_VAR_HEAD
    long ob_shash; // init as -1
    int ob_sstate; // whether interned
    char ob_sval[1]; //trick, 'cos this is not only for gnu c
} PyStringObject;
```


stringobject.h

Built-in Types/String

- not changable.
 - the good: used as key in dict
 - the bad: efficiency in join
- you can have `\0` in your string, for size is determined by `ob_size` and `tp_itemsize`
- `PyString_FromString` `PyString_FromStringAndSize`
- intern

Built-in Types/String

字符串"Python"所占内存示意图



Built-in Types/String

intern

- a cache mechanism, working with hash can speed up python by 20%
- happens at compile time, not run time
- by default, only cache strings consist of numbers, alphabets and underscores, but built-in `intern()` can be used to cache any string.
- internally, use a dict `<PyObject *, PyObject *>` named 'interned' to store, and the two refs in interned don't count.
- nullstring and characters array use intern
- names like `"__name__"`, `"__doc__"` are interned
- see implementations in `codeobject.c` `stringobject.c`

Built-in Types/String

efficiency of `"+="`/`"+"` and `"join"`

- `string_join` allocate memory once for all the items
- `+=` and `+` allocate memory everytime it executes, but optimized for short string, which resize the lvalue. see `BINARY_ADD`, `INPLACE_ADD`, `string_concatenate` in `Python/ceval.c`
- older version of concatenate is `string_concat` in `Object/stringobject.c` which malloc a new memory

PEP8: do not rely on CPython's efficient implementation of in-place string concatenation for statements in the form `a += b` or `a = a + b`. This optimization is fragile even in CPython (it only works for some types) and isn't present at all in implementations that don't use refcounting.

Built-in Types/String

stringlib fastsearch

- based on a mix between boyer-moore and horspool
- 'in' operation

Built-in Types/List

```
typedef struct {
    PyObject_VAR_HEAD
    /* Vector of pointers to list elements. list[0] is ob_item
       [0], etc. */
    PyObject **ob_item;
    /* ob_item contains space for 'allocated' elements. The
       number * currently in use is ob_size.
       * Invariants:
       *     0 <= ob_size <= allocated
       *     len(list) == ob_size
       *     ob_item == NULL implies ob_size == allocated == 0
       * list.sort() temporarily sets allocated to -1 to detect
         mutations.
       * Items must normally not be NULL, except during
         construction when * the list is not yet visible
         outside the function that builds it. */
    Py_ssize_t allocated;
} PyObject;
```

Built-in Types/List

- `ob_size` and `allocated` is like `size` and `capacity` of `vector` in `c++`
- only one method to init: `PyList_New`
- cache mechanism: `free_list`, which hold 80 objects
- `list_dealloc` to return `PyObject` to `free_list`
- memory management in `list_resize`:
 - `newsize < allocated && newsize > allocated/2`, no need to `realloc`
 - (0, 4, 8, 16, 25, 35, 46, 58, 72, 88)
 - otherwise shrink or enlarge

Built-in Types/List

misc notes

- insert when $> \text{len}(\text{list})$ or $< -\text{len}(\text{list})$
- xrange / range
- tim sort
- consider array
- consider tuple
 - allocate once
 - cache 2000, each < 20 elements
 - read-only, better fit for concurrency

Built-in Types/Dict

```
typedef struct _dictobject PyDictObject;
struct _dictobject {
    PyObject_HEAD
    Py_ssize_t ma_fill; /* # Active + # Dummy */
    Py_ssize_t ma_used; /* # Active */
    Py_ssize_t ma_mask;
    PyDictEntry *ma_table;
    PyDictEntry *(*ma_lookup)(PyDictObject *mp, PyObject *key,
        long hash);
    PyDictEntry ma_smalltable[PyDict_MINSIZE];
};
```

dictobject.h

Built-in Types/Dict

```
typedef struct {  
    /* Cached hash code of me_key.  Note that hash codes are C  
       longs.  
     * We have to use Py_ssize_t instead because dict_popitem()  
       abuses  
     * me_hash to hold a search finger.  
     */  
    Py_ssize_t me_hash;  
    PyObject *me_key;  
    PyObject *me_value;  
} PyDictEntry;
```

dictobject.h

Built-in Types/Dict

- hash table with open addressing
- unused/active/dummy
- ma_smalltable and ma_table
- lookdict and lookdict_string
- PyDict_New
- cache mechanism: free_list, which hold 80 objects
- memory management in dictresize:
 - if fill $> 2/3$ size, adjust size.
 - size * 2 or size * 4

Memory Allocators

```
30  Object-specific allocators
31  _____
32  [ int ] [ dict ] [ list ] ... [ string ]      Python core      |
33 +3 | <----- Object-specific memory -----> | <-- Non-object memory --> |
34  |                                           |                   |
35  [ Python's object allocator ]                |                   |
36 +2 | ##### Object memory ##### | <----- Internal buffers -----> |
37  |                                           |                   |
38  [ Python's raw memory allocator (PyMem_ API) ] |                   |
39 +1 | <----- Python memory (under PyMem manager's control) -----> | |
40  |                                           |                   |
41  [ Underlying general-purpose allocator (ex: C library malloc) ] |
42 0 | <----- Virtual memory allocated for the python process -----> |
43
44  =====
```


Memory Allocators

- Layer 0 is OS's allocators
- Layer 1 is just a simple wrapper of Layer 0, in forms of both macro and function
- Layer 2 is the core mechanism—pymalloc.
PyObject_{Malloc,Realloc,Free}
- Layer 3 is type specific memory cache

- block: works for $<256B$, when $>256B$, use Layer 1. block must be 8B alignment.
- pool: same as mem page size(4KB), manage blocks of same size
- arena: has 64 pools.
- usedpools: array of head pointers of USED pool

Mark-Sweep Garbage Collection

- break unreachable reference cycles
- doesn't explicitly deallocate any objects, just breaks cycles
- automatically run by interpreter every once in a while.
`sys.getcheckinterval()`
- auto gc object without `__del__`
- stop the world to gc
- expensive, cost is linear to number of references of vector in program.

Generational GC

- most objects live either for a very short time or for a very long time
- divide object into 3 generations(0, 1, 2)
- each new object is generation 0.
- n-th gen. gc analyses objects of gen. 0-n
- any object that survived n-th gen. gc is promoted to n+1-th gen.
- (700, 10, 10) 700 is alloc - dealloc, 10s are times of prior gen gc.

Memory Leak

- leak in modules written in c/c++
- integers and floats
- gc won't collect any objects in cycles where at least one object has `__del__`
- hiding references

Memory Leak

Hidding References

- closures, `functools.partial`, etc
- `sys.exc_traceback` keeps last exception handled in this stack frame including the whole stack state when exception occurs.
- `sys.last_traceback` keeps unhandled excption including whole stack state

What to do

- use context manager("with") instead of `__del__`
- use `weakref`
- if you really need `__del__`, be careful of ref cycle
- tools: `pdb/objgraph/guppy/heapy/memory_profiler/line_profiler/gc.collect()`

Q & A